


Computation & ML models





Overview

- Attendance
- Computational complexity
- Network data structures
- Metric algorithms
- Software



Computational complexity





Computational complexity

- A measure of a computational algorithm's running time
- For graph algorithms, this will usually be a function of n (number of nodes) and/or m (number of edges)
- We will use Big O notation where $O(f(x))$ implies the run time is proportional to $f(x)$ to leading order.
- Examples:
 - $O(1)$ constant run time
 - $O(n)$ run time proportional to n
 - $O(n^2)$ run time proportional to n^2




Image credit: wikipedia.org



Will my algorithm finish while I get coffee?

- Consider an algorithm with complexity $O(f(x))$. For an arbitrary value x_1 , the run time is $t_1 = \beta f(x_1) \rightarrow \beta = t_1 / f(x_1)$.
 - Consider $x_2 = \alpha x_1$, then
$$t_2 = \beta f(x_2) = \left[f(\alpha x_1) / f(x_1) \right] t_1$$
 - If $f(x) \equiv x^k$, then $f(\alpha x) \equiv \alpha^k x^k$ and $t_2 = \alpha^k t_1$
- Suppose we have an algorithm (or a machine learning training step) for a graph with complexity $O(m + n)$ (linear in $m + n$)
 - Test run on a graph with $n = 10^3$ and $m = 10^4$ takes 1 second
 - For a network 1000 times larger ($n = 10^6$ and $m = 10^7$) runtime will be about 1000 seconds (16.67 minutes) (note here $\alpha = 1000$ and $k = 1$)
- What if the algorithm complexity is $O(n^3)$?
 - Consider the same test run configuration with 1 second running time
 - Scaling the graph by 1000 ($n = 10^6$ and $m = 10^7$), implies a run time of 1 billion seconds (**31.7 years**) (note here $\alpha = 1000$ and $k = 3$)



More on computational complexity

- For graph networks, any algorithm with complexity $O(n^3)$ or more is too slow for large networks
- The time estimates for an algorithm for a scaled-up network based on a test run and algorithm complexity assume the scaled network fits in RAM
 - If not, the algorithm time on the larger network is likely to be much longer than suggested by the scaling argument



Network data structures





How should we represent the network on a computer?


- The data structure (in RAM) used to represent the network can have a significant effect on algorithm performance
- In all representations, we'll label nodes with integers
 - In the math, we'll use $1, \dots, n$
 - In Python, we'll use $0, \dots, n - 1$
 - Most graph software will handle this for us, though in our deep learning models we may need to handle this manually
- Nodes often have attributes. These can be stored in an appropriate data structure with index (or key) access (e.g., list of node names)

Adjacency matrix representation

- Label the nodes with integers $1 \cdots n$ (or $0, \cdots, n - 1$)
- For a simple, undirected network, the **adjacency matrix**, A , is the $n \times n$ matrix with elements a_{ij}

$$a_{ij} = \begin{cases} 1 & \text{if there is an edge between nodes } i \text{ and } j \\ 0 & \text{otherwise} \end{cases}$$

- We can represent this in computer memory using any 2D array data structure (Python lists, NumPy array, PyTorch tensor)



```

[In [26]: A = [[0,1,1,1],[1,0,0,1],[1,0,0,0],[1,1,0,0]]

[In [27]: Anp = np.array(A)

[In [28]: Anp
Out[28]:
array([[0, 1, 1, 1],
       [1, 0, 0, 1],
       [1, 0, 0, 0],
       [1, 1, 0, 0]])


[In [29]: Apt = torch.tensor(A, dtype=torch.int32)

[In [30]: Apt
Out[30]:
tensor([[0, 1, 1, 1],
        [1, 0, 0, 1],
        [1, 0, 0, 0],
        [1, 1, 0, 0]], dtype=torch.int32)
  
```



Adjacency matrix representation disadvantages


- Finding the neighbors of a given node is $O(n)$
 - We must iterate over the corresponding row in the array
 - Many algorithms require doing this repeatedly
- For sparse networks, most entries are 0
 - Inefficient use of computer memory
 - Assuming 4 byte integers, then a network with n nodes requires $4n^2$ bytes to store the adjacency matrix
 - 50,000 nodes requires 10 GB of RAM
 - 10^6 nodes requires 4000 GB of RAM – this network size is not uncommon





Adjacency list

- Store a list for each node where the list for node v contains the labels of its neighbors
- Each edge appears twice in the adjacency list




Node	Neighbors
1	2, 3
2	1, 4
3	1
4	1, 2



Adjacency list

- Store a list for each node where the list for node i contains the labels of its neighbors
- Each edge appears twice in the adjacency list
- Special considerations
 - Multiedges are represented by repeated entries in the neighbors list




Node	Neighbors
1	2, 3, 3, 4
2	1, 4
3	1, 1
4	1, 2



Adjacency list

- Store a list for each node where the list for node i contains the labels of its neighbors
- Each edge appears twice in the adjacency list
- Special considerations
 - Multiedges are represented by repeated entries in the neighbors list
 - Self-loop (node connected to itself) is represented by the node being included in its neighbor list (**included twice**)




Node	Neighbors
1	2, 3, 4
2	1, 4, 2, 2
3	1
4	1, 2



Adjacency list

- Store a list for each node where the list for node i contains the labels of its neighbors
- Each edge appears twice in the adjacency list
- Special considerations
 - Multiedges are represented by repeated entries in the neighbors list
 - Self-loop (node connected to itself) is represented by the node being included in its neighbor list (**included twice**)
- Directed network can be represented by two lists, one for outgoing edges, one for incoming edges



Node	Out	In
1	2, 3	4
2		1, 4
3		1
4	2, 4	4



Adjacency list advantages


- Only requires $2m$ integers to represent the network, i.e., the memory for storage scales linearly with m
 - A network with $m = 10^5$ edges and $n = 10^4$ nodes requires 800 kB compared to 400 MB for the adjacency matrix
 - A network with $m = 10^7$ edges and $n = 10^6$ nodes requires 80 MB compared to 4000 GB for the adjacency matrix
- Enumeration of node neighbors is $O(m/n)$ (average) compared to $O(n)$ for adjacency matrix



Adjacency lists in Python

We can use nested lists or a hash (dictionary)

```
AL = [[1,2],[0,3],[0],[0,1]]
ALd = {0:[1,2], 1:[0,3], 2:[0], 3:[0,1]}
```



We cannot use numpy arrays or torch tensors. Why?

```

[In [11]: np.array(AL)
-----
ValueError                                Traceback (most recent call last)
Cell In[11], line 1
----> 1 np.array(AL)

ValueError: setting an array element with a sequence. The requested array has an inhomogeneous shape after 1 dimensions. The detected shape was (4,) + inhomogeneous part.

[In [12]: torch.tensor(AL)
-----
ValueError                                Traceback (most recent call last)
Cell In[12], line 1
----> 1 torch.tensor(AL)

ValueError: expected sequence of length 2 at dim 1 (got 1)
  
```

Node	Neighbors
0	1, 2
1	0, 3
2	0
3	0, 1

We can instead use `scipy.sparse` and `torch.sparse` classes (more on this later)




Metric algorithms





Node degree from adjacency matrix

- Node degree is the number of edges connected to the node
- The algorithm (and complexity) depend on the data structure
- Consider node with index i
- Adjacency matrix representation algorithm
 1. Access row i
 2. Iterate over the row elements and count non-zero entries (sum for unweighted networks)




	0	1	0	0	0	0	0	1
0	1	0	1	0	0	0	0	0
1	0	1	0	0	0	0	0	1
2	0	0	0	0	0	0	0	1
3	0	0	0	0	0	1	0	0
4	0	0	0	0	1	0	1	0
5	0	0	0	0	0	1	0	1
6	1	0	1	1	0	0	1	0

Node **degree for node 0 is 2**, which is the sum of row zero in the adjacency matrix



Node degree from adjacency matrix

- Node degree is the number of edges connected to the node
- The algorithm (and complexity) depend on the data structure
- Consider node with index i
- Adjacency matrix representation algorithm
 1. Access row i which
 2. Iterate over the row elements and count non-zero entries (sum for unweighted networks)




0	1	0	0	0	0	0	1
1	0	1	0	0	0	0	0
0	1	0	0	0	0	0	1
0	0	0	0	0	0	0	1
0	0	0	0	0	1	0	0
0	0	0	0	1	0	1	0
0	0	0	0	0	1	0	1
1	0	1	1	0	0	1	0

Node **degree for node 3 is 1**, which is the sum of row zero in the adjacency matrix



Node degree from adjacency matrix

- Node degree is the number of edges connected to the node
- The algorithm (and complexity) depend on the data structure
- Consider node with index i
- Adjacency matrix representation algorithm
 1. Access row i which
 2. Iterate over the row elements and count non-zero entries (sum for unweighted networks)




0	1	0	0	0	0	0	1
1	0	1	0	0	0	0	0
0	1	0	0	0	0	0	1
0	0	0	0	0	0	0	1
0	0	0	0	0	1	0	0
0	0	0	0	1	0	1	0
0	0	0	0	0	1	0	1
1	0	1	1	0	0	1	0

Node **degree for node 7 is 4**, which is the sum of row zero in the adjacency matrix

Node degree from adjacency list

- Node degree is the number of edges connected to the node
- The algorithm (and complexity) depend on the data structure
- Consider node with index i
- Adjacency list representation
 1. Access i^{th} list (or dict) entry
 2. Obtain list length (e.g., `len (AL (i) )` in Python)



0: 1, 7

1: 0, 2

2: 1, 7

3: 7

4: 5

5: 4, 6

6: 5, 7


7: 0, 2, 3, 6

Node **degree for node 0 is 2**, which is the length of the list for node 0

*Python (like other modern programming languages) stores arrays as contiguous blocks in memory

Node degree from adjacency list

- Node degree is the number of edges connected to the node
- The algorithm (and complexity) depend on the data structure
- Consider node with index i
- Adjacency list representation
 1. Access i^{th} list (or dict) entry
 2. Obtain list length (e.g., `len (AL (i) )` in Python)



0: 1, 7

1: 0, 2

2: 1, 7

3: 7

4: 5

5: 4, 6

6: 5, 7


7: 0, 2, 3, 6

Node **degree for node 3 is 1**, which is the length of the list for node 0

*Python (like other modern programming languages) stores arrays as contiguous blocks in memory

Node degree from adjacency list

- Node degree is the number of edges connected to the node
- The algorithm (and complexity) depend on the data structure
- Consider node with index i
- Adjacency list representation
 1. Access i^{th} list (or dict) entry
 2. Obtain list length (e.g., `len (AL (i) )` in Python)



0: 1, 7

1: 0, 2

2: 1, 7

3: 7

4: 5

5: 4, 6

6: 5, 7

7: 0, 2, 3, 6


Node **degree for node 7 is 4**, which is the length of the list for node 0

*Python (like other modern programming languages) stores arrays as contiguous blocks in memory



Node degree complexity is affected by representation

- Adjacency matrix representation algorithm
 - *Access row i which is $O(1)$
 - Iterate over the row elements and count non-zero entries (sum for unweighted networks) which is $O(n)$
- Adjacency list representation
 - Access i^{th} list (or dict) entry which is $O(1)$
 - Obtain list length (e.g., `len (AL (i) )` in Python) which is $O(1)$



https://github.com/masino-teaching/ml4g/blob/main/notebooks/network_node_degree_complexity.ipynb

*Python (like other modern programming languages) stores arrays as contiguous blocks in memory




Local clustering coefficient computation

$$C = \frac{\text{Number of pairs of neighbors of } i \text{ that are connected}}{\text{number pairs of neighbors of } i}$$

- Check every distinct pair of neighbors of i of which there are $\binom{k_i}{2} = \frac{k_i(k_i-1)}{2}$ to see if they're connected
- We only need to check the pair once, so restrict our check to neighbors u, v such that $u < v$
- Utilize an adjacency list representation

Image credit: Newman 2018






Local clustering coefficient computation

```
// Input: Graph G, target node i
neighbors_i ← GetNeighbors(G, i)
connected_pairs ← 0

FOR each neighbor u in neighbors_i DO
  FOR each neighbor v in neighbors_i DO
    IF u < v THEN
      IF EdgeExists(G, u, v) THEN
        connected_pairs ← connected_pairs + 1
      END IF
    END IF
  END FOR
END FOR

k ← |neighbors_i|
C_i ← 2 × connected_pairs / (k × (k-1))
RETURN C_i
```



0: [1, 2]
1: [0, 2]
2: [0, 1, 3, 4]
3: [2, 4]
4: [2, 3]

Step	u	v	Check Edge	Edge Exists?	connected_pairs	Action
1	0	1	$1 \in \text{adj}[0]$	YES	1	+1

<https://ml4g.masinolab.com/topics/metric-algorithms/local-clustering-coefficient-computation.html>




Local clustering coefficient computation

```
// Input: Graph G, target node i
neighbors_i ← GetNeighbors(G, i)
connected_pairs ← 0

FOR each neighbor u in neighbors_i DO
  FOR each neighbor v in neighbors_i DO
    IF u < v THEN
      IF EdgeExists(G, u, v) THEN
        connected_pairs ← connected_pairs + 1
      END IF
    END IF
  END FOR
END FOR

k ← |neighbors_i|
C_i ← 2 × connected_pairs / (k × (k-1))
RETURN C_i
```



```
0: [1, 2]
1: [0, 2]
2: [0, 1, 3, 4]
3: [2, 4]
4: [2, 3]
```

Step	u	v	Check Edge	Edge Exists?	connected_pairs	Action
1	0	1	$1 \in \text{adj}[0]$	YES	1	+1
2	0	3	$3 \in \text{adj}[0]$	NO	1	skip

<https://ml4g.masinolab.com/topics/metric-algorithms/local-clustering-coefficient-computation.html>




Local clustering coefficient computation

```
// Input: Graph G, target node i
neighbors_i ← GetNeighbors(G, i)
connected_pairs ← 0

FOR each neighbor u in neighbors_i DO
  FOR each neighbor v in neighbors_i DO
    IF u < v THEN
      IF EdgeExists(G, u, v) THEN
        connected_pairs ← connected_pairs + 1
      END IF
    END IF
  END FOR
END FOR

k ← |neighbors_i|
C_i ← 2 × connected_pairs / (k × (k-1))
RETURN C_i
```



0: [1, 2]
1: [0, 2]
2: [0, 1, 3, 4]
3: [2, 4]
4: [2, 3]

Step	u	v	Check Edge	Edge Exists?	connected_pairs	Action
1	0	1	$1 \in \text{adj}[0]$	YES	1	+1
2	0	3	$3 \in \text{adj}[0]$	NO	1	skip
3	0	4	$4 \in \text{adj}[0]$	NO	1	skip

<https://ml4g.masinolab.com/topics/metric-algorithms/local-clustering-coefficient-computation.html>




Local clustering coefficient computation

```
// Input: Graph G, target node i
neighbors_i ← GetNeighbors(G, i)
connected_pairs ← 0

FOR each neighbor u in neighbors_i DO
  FOR each neighbor v in neighbors_i DO
    IF u < v THEN
      IF EdgeExists(G, u, v) THEN
        connected_pairs ← connected_pairs + 1
      END IF
    END IF
  END FOR
END FOR

k ← |neighbors_i|
C_i ← 2 × connected_pairs / (k × (k-1))
RETURN C_i
```



```
0: [1, 2]
1: [0, 2]
2: [0, 1, 3, 4]
3: [2, 4]
4: [2, 3]
```

Step	u	v	Check Edge	Edge Exists?	connected_pairs	Action
1	0	1	$1 \in \text{adj}[0]$	YES	1	+1
2	0	3	$3 \in \text{adj}[0]$	NO	1	skip
3	0	4	$4 \in \text{adj}[0]$	NO	1	skip
4	1	3	$3 \in \text{adj}[1]$	NO	1	skip

<https://ml4g.masinolab.com/topics/metric-algorithms/local-clustering-coefficient-computation.html>




Local clustering coefficient computation

```
// Input: Graph G, target node i
neighbors_i ← GetNeighbors(G, i)
connected_pairs ← 0

FOR each neighbor u in neighbors_i DO
  FOR each neighbor v in neighbors_i DO
    IF u < v THEN
      IF EdgeExists(G, u, v) THEN
        connected_pairs ← connected_pairs + 1
      END IF
    END IF
  END FOR
END FOR

k ← |neighbors_i|
C_i ← 2 × connected_pairs / (k × (k-1))
RETURN C_i
```



```
0: [1, 2]
1: [0, 2]
2: [0, 1, 3, 4]
3: [2, 4]
4: [2, 3]
```

Step	u	v	Check Edge	Edge Exists?	connected_pairs	Action
1	0	1	1 ∈ adj[0]	YES	1	+1
2	0	3	3 ∈ adj[0]	NO	1	skip
3	0	4	4 ∈ adj[0]	NO	1	skip
4	1	3	3 ∈ adj[1]	NO	1	skip
5	1	4	4 ∈ adj[1]	NO	1	skip
6	3	4	4 ∈ adj[3]	YES	2	+1

<https://ml4g.masinolab.com/topics/metric-algorithms/local-clustering-coefficient-computation.html>

$$C = \frac{2 \times \text{connected_pairs}}{k_i(k_i - 1)} = \frac{2 \times 2}{4 \times 3} \approx 0.333$$



Local clustering coefficient computation complexity

- The two outer loops combined have approximately $O(k_i^2)$
- The complexity of checking if an edge exists between nodes u and v depends on the node structure
 - For a network with degree homophily $O(k_i)$
 - For a network without degree homophily $O\left(\frac{2m}{n}\right)$
- The overall algorithm complexity is
 - For a network with degree homophily $O(k_i^3)$
 - For a network without degree homophily $O\left(k_i^2 \frac{2m}{n}\right)$
- For “scale-free” networks (i.e., those with a power law distribution on node degree), the complexity is effectively unbounded because k_i is very large.

```
// Input: Graph G, target node i
neighbors_i ← GetNeighbors(G, i)
connected_pairs ← 0

FOR each neighbor u in neighbors_i DO
  FOR each neighbor v in neighbors_i DO
    IF u < v THEN
      IF EdgeExists(G, u, v) THEN
        connected_pairs ← connected_pairs + 1
      END IF
    END IF
  END FOR
END FOR

k ← |neighbors_i|
C_i ← 2 × connected_pairs / (k × (k-1))
RETURN C_i
```



Breadth-first search

- Finds the **shortest distance** from a starting node, s , to all other nodes
- We will need this to compute other metrics
- General concept
 1. Set $d = 0$. Set distance to s to d , all other nodes set distance to undefined.
 2. Find all nodes with current distance, d .
 - If there are no such nodes, stop
 - In the first round, s , will be the only node at distance $d = 0$
 3. For each node, i , at the current distance, d , check its neighbors
 - If the neighbor has a previously assigned distance, do nothing
 - Else, assign it distance $d + 1$
 - Increment d , Go back to step 2




Image credit: Newman 2018



Breadth-first search implementation

- Naïve approach
 1. Create an array, *dis*, of n integers to store distance of each node from s
 2. Set the distance from s to itself to 0, all others to -1
 3. Set $d = 0$
 4. **Examine the *dis* array to find all nodes that are distance d from s**
 5. Find the neighbors of those nodes and set them to $d + 1$ if their value is -1 (otherwise leave as is)
 6. Increase d by 1
 7. Repeat from step 4

- In the worst case, this algorithm has $O(m + n^2)$ due to the repeated checks of the distance array
- Observe that in each round, the nodes whose distances we set to $d + 1$ are the ones that are distance d from s in the next round
- If we store these in a queue, we can get a much more efficient algorithm with $O(m + n)$ complexity

https://ml4g.masinolab.com/topics/metric-algorithms/bfs_algorithm_visualization.html




Finding shortest paths algorithm

- Goal is to find ALL shortest paths from a node s to any other node
- There may be multiple shortest paths


Pseudo Algorithm:

1. **Create dictionary 'paths':** Keys are node labels, values are lists of predecessor nodes
 2. **Conduct BFS:** When examining neighbor j of node i at distance d :
 - **Case i:** If j is previously unseen $\rightarrow$ set $\text{distance}(j) = d+1$, add i to $\text{paths}[j]$
 - **Case ii:** If j is previously seen and has distance $d+1 \rightarrow$ add i to $\text{paths}[j]$
 - **Case iii:** If j has distance $< d+1 \rightarrow$ do nothing
 3. **Reconstruct paths:** Starting at any node, follow predecessor chains in $\text{paths}[]$ back to source s
- Complexity is the same as our BFS algorithm as the only extra step is storing the paths dictionary which has $O(1)$ complexity



There are two *shortest paths* from node 0 to node 4

Finding shortest paths algorithm



<https://ml4g.masinolab.com/topics/metric-algorithms/finding-shortest-paths-algorithm.html>

Step: 1

Status: Processing node $i=0$ at distance $d=0$

Queue: [1, 2]

Current node i : 0

Neighbors of node i : [1, 2]

Actions for each neighbor:

Neighbor 1 unseen. Setting distance to 1. Adding 0 to paths[1].
Neighbor 2 unseen. Setting distance to 1. Adding 0 to paths[2].


Distances:

0: 0
1: 1
2: 1

Paths (Predecessors):

0: []
1: [0]
2: [0]

Finding shortest paths algorithm



<https://ml4g.masinolab.com/topics/metric-algorithms/finding-shortest-paths-algorithm.html>

Step: 2

Status: Processing node $i=1$ at distance $d=1$

Queue: [2, 3, 4]

Current node i : 1

Neighbors of node i : [0, 2, 3, 4]

Actions for each neighbor:

Neighbor 0 has shorter distance $0 < 2$. Do nothing.

Neighbor 2 has shorter distance $1 < 2$. Do nothing.

Neighbor 3 unseen. Setting distance to 2. Adding 1 to paths[3].

Neighbor 4 unseen. Setting distance to 2. Adding 1 to paths[4].


Distances:

```
0: 0
1: 1
2: 1
3: 2
4: 2
```

Paths (Predecessors):

```
0: []
1: [0]
2: [0]
3: [1]
4: [1]
```

Finding shortest paths algorithm



<https://ml4g.masinolab.com/topics/metric-algorithms/finding-shortest-paths-algorithm.html>

Step: 3

Status: Processing node $i=2$ at distance $d=1$

Queue: [3, 4]

Current node i : 2

Neighbors of node i : [0, 1, 4]

Actions for each neighbor:

Neighbor 0 has shorter distance $0 < 2$. Do nothing.

Neighbor 1 has shorter distance $1 < 2$. Do nothing.

Neighbor 4 already at distance 2. Adding 2 to paths[4].

Distances:

```

0: 0
1: 1
2: 1
3: 2
4: 2
  
```


Paths (Predecessors):

```

0: []
1: [0]
2: [0]
3: [1]
4: [1, 2]
  
```



Finding shortest paths algorithm



All Shortest Paths from 0

To node 0:

0

To node 1:

0 → 1

To node 2:

0 → 2

To node 3:

0 → 1 → 3

To node 4:

0 → 1 → 4

0 → 2 → 4

To node 5:

0 → 1 → 3 → 5

0 → 1 → 4 → 5

0 → 2 → 4 → 5

To node 6:

0 → 1 → 4 → 6

0 → 2 → 4 → 6

Step: 7

Status: Algorithm complete - reconstructing paths

Queue: []

Current node i: 6

Neighbors of node i: [4, 5]

Actions for each neighbor:

Neighbor 4 has shorter distance $2 < 4$. Do nothing.

Neighbor 5 has shorter distance $3 < 4$. Do nothing.

Distances:

0: 0

1: 1

2: 1

3: 2

4: 2

5: 3

6: 3

Paths (Predecessors):

0: []

1: [0]

2: [0]

3: [1]

4: [1, 2]

5: [3, 4]

6: [4]

<https://ml4g.masinolab.com/topics/metric-algorithms/finding-shortest-paths-algorithm.html>



Betweenness Centrality

- Measures the extent to which a node is on shortest paths between other nodes
- The betweenness centrality, x_i , of node i is

$$x_i = \sum_{st} \frac{n_{st}^i}{g_{st}}$$

where the sum is over all shortest paths for all node pairs $\{s, t\}$, n_{st}^i is the number of shortest paths from s to t that contain i , and g_{st} is the number of shortest paths from s to t

- Naively, we could just find all the shortest paths between s and t and check them all for node i .
 - Implementing this for all pairs s and t has complexity $O(n^{5/2})$



Betweenness Centrality Algorithm – part 1

- Modify the BFS + shortest paths algorithm to store the number of shortest paths that pass through a node.
 - Consider an arbitrary source node s
1. Initialize a distance array $\mathbf{d}$, a weight array, $\mathbf{w}$, and a queue array, $\mathbf{q}$, each of length n all with -1 entries. Initialize a *paths* dictionary to store the shortest paths. Initialize a read pointer, $p_r = 0$ and a write pointer, $p_w = 1$.
 2. Set $\mathbf{d}[s] = 0$. Set $\mathbf{w}[s] = 1$. Set $\mathbf{q}[p_r] = s$.
 3. If $p_r = p_w$, stop, else obtain node i from $\mathbf{q}[p_r]$ and set $p_r = p_r + 1$. (The first time, $\mathbf{q}[p_r] = s$).
 4. For each neighbor j of i :
 - a. If $\mathbf{d}[j] == -1$, set $\mathbf{d}[j] = \mathbf{d}[i] + 1$ and set $\mathbf{w}[j] = \mathbf{w}[i]$. Set $\text{paths}[j] = [i]$ Set $\mathbf{q}[p_w] = j$. Set $p_w = p_w + 1$.
 - b. If $\mathbf{d}[j] == \mathbf{d}[i] + 1$, set $\mathbf{w}[j] = \mathbf{w}[j] + \mathbf{w}[i]$ and set $\text{paths}[j] = \text{paths}[j] + [i]$
 - c. Otherwise do nothing
 5. Repeat from step 3

<https://ml4g.masinolab.com/topics/metric-algorithms/betweenness-centrality-algorithm.html>



Betweenness Centrality Algorithm – part 2

1. Find every “leaf” node, t . This can be done by finding the nodes that do not appear in a value of the *paths* dictionary (i.e., these nodes are not on any shortest path). Store these in a set, *leaves*.
2. Initialize a score array $\mathbf{x}$ of size n . For each leaf node, l , in *leaves* set $\mathbf{x}(l) = 1$.
3. For each node i in *reversed*($\mathbf{q}$):
 - a. If i in *leaves*, continue to next value of i
 - b. Initialize an empty list, *nb* to store neighbors that precede i on a shortest path to s
 - c. For node j in all nodes:
 - i. If i in *paths*[j], add j to *nb*
 - d. Set $\mathbf{x}[i] = 1 + \sum_{j \in nb} \frac{x_j w_{ji}}{w_j}$

<https://ml4g.masinolab.com/topics/metric-algorithms/betweenness-centrality-algorithm.html>



Betweenness centrality complexity

- The algorithm in the previous slides computes the betweenness scores for all nodes for a single source node, s
- This must be repeated for all possible sources nodes
- Overall complexity is
 - $O(n(n + m))$ in general
 - $O(n^2)$ for sparse networks

<https://ml4g.masinolab.com/topics/metric-algorithms/betweenness-centrality-algorithm.html>



Software





Software in this course

We will use a Python software stack that includes:

- PyTorch – general deep learning library
- PyTorch Geometric – extension of PyTorch for graph neural networks
- NetworkX – network analysis and visualization
- PyGraphviz – Python interface to Graphviz network visualization library

We will work with these starting in lab 1 (next meeting)

Please review the Google-Colab-Palmetto.docx file on Canvas!!!!



Other software

For the intrepid learners (we won't cover all of these)

- Deep Graph Library (dgl.ai)

Name	Availability	Platform	Description
Gephi	Free	WML	Interactive network analysis and visualization
Pajek	Free	W	Interactive social network analysis and visualization
InFlow	Commercial	W	Interactive social network analysis and visualization
UCINET	Commercial	W	Interactive social network analysis
yEd	Free	WML	Interactive visualization
Visone	Free	WL	Interactive visualization
Graphviz	Free	WML	Visualization
NetworkX	Free	WML	Python library for network analysis and visualization
JUNG	Free	WML	Java library for network analysis and visualization
igraph	Free	WML	C/R/Python libraries for network analysis

PyTorch Tutorials

<https://pytorch.org/tutorials/>

 PyTorch

[Get Started](#) [Ecosystem](#) [Edge](#) [Blog](#) [Tutorials](#) [Docs](#) [Resources](#)

2.2.1+cu121

 Search Tutorials

[PyTorch Recipes](#) [Introduction to PyTorch](#)

[Learn the Basics](#)
[Quickstart](#)
[Tensors](#)
[Datasets & DataLoaders](#)
[Transforms](#)
[Build the Neural Network](#)
[Automatic Differentiation with torch.autograd](#)
[Optimizing Model Parameters](#)
[Save and Load the Model](#)

[Introduction to PyTorch on YouTube](#)

[Introduction to PyTorch - YouTube Series](#)
[Introduction to PyTorch](#)
[Introduction to PyTorch Tensors](#)
[The Fundamentals of Autograd](#)
[Building Models with PyTorch](#)
[PyTorch TensorBoard Support](#)
[Training with PyTorch](#)
[Model Understanding with Captum](#)

Tutorials > Welcome to PyTorch Tutorials

Welcome to PyTorch Tutorials

What's new in PyTorch tutorials?

- [PyTorch Inference Performance Tuning on AWS Graviton Processors](#)
- [Using TORCH_LOGS python API with torch.compile](#)
- [PyTorch 2 Export Quantization with X86 Backend through Inductor](#)
- [Getting Started with DeviceMesh](#)
- [Compiling the optimizer with torch.compile](#)

Learn the Basics

Familiarize yourself with PyTorch concepts and modules. Learn how to load data, build deep neural networks, train and save your models in this quickstart guide.

[Get started with PyTorch](#)

PyTorch Recipes

Bite-size, ready-to-deploy PyTorch code examples.

[Explore Recipes](#)

AllAttentionAudioAxBackendsBest PracticeC++CUDA

Extending PyTorchFXFrontend APIsGetting StartedImage/Video

https://pytorch-geometric.readthedocs.io/en/stable/get_started/introduction.html

Introduction by Example

- Data Handling of Graphs
- Common Benchmark Datasets
- Mini-batches
- Data Transforms
- Learning Methods on Graphs
- Exercises

Colab Notebooks and Video Tutorials

TUTORIALS

- Design of Graph Neural Networks
- Working with Graph Datasets
- Use-Cases & Applications
- Distributed Training

ADVANCED CONCEPTS

- Advanced Mini-Batching
- Memory-Efficient Aggregations
- Hierarchical Neighborhood Sampling
- Compiled Graph Neural Networks
- TorchScript Support
- Scaling Up GNNs via Remote Backends
- Managing Experiments with GraphGym
- CPU Affinity for PyG Workloads

PACKAGE REFERENCE

- `torch_geometric`
- `torch_geometric.nn`

🏠 / Introduction by Example

Introduction by Example

We shortly introduce the fundamental concepts of 🧠 PyG through self-contained examples.

For an introduction to Graph Machine Learning, we refer the interested reader to the 📖 Stanford CS224W: Machine Learning with Graphs lectures. For an interactive introduction to 🧠 PyG, we recommend our carefully curated 📄 Google Colab notebooks.

At its core, 🧠 PyG provides the following main features:

- Data Handling of Graphs
- Common Benchmark Datasets
- Mini-batches
- Data Transforms
- Learning Methods on Graphs
- Exercises

Data Handling of Graphs

A graph is used to model pairwise relations (edges) between objects (nodes). A single graph in 🧠 PyG is described by an instance of `torch_geometric.data.Data`, which holds the following attributes by default:

- `data.x` : Node feature matrix with shape `[num_nodes, num_node_features]`
- `data.edge_index` : Graph connectivity in COO format with shape `[2, num_edges]` and type `torch.long`
- `data.edge_attr` : Edge feature matrix with shape `[num_edges, num_edge_features]`

