


Shallow Node Embeddings





Overview

- Motivation
- Shallow node embedding framework
- Node embedding methods
 - Factorization approaches
 - Random walk approaches
- Graph embedding
- Limitations of matrix factorization and random walk embeddings



Motivation






Previous lab – Machine learning with engineered features

- In a standard machine learning approach, we extract mathematically formulated metrics from an input graph to train and make inferences with a statistical learning model
- As with many other areas (e.g., computer vision, natural language processing), it is frequently beneficial to automate feature learning and bypass feature engineering

Input graph



Engineered Features (centrality, cliques, k-cores, diameter, etc)


Machine learning algorithm

Inference



Latent representation learning for graphs

We would like to learn node, edge, and graph representations (embeddings) that are useful for many tasks



$$f: u \rightarrow \mathbb{R}^d$$



Vector in $\mathbb{R}^d$



- f is a function that maps an input, u , (node, edge, or graph) to numerical vector in $\mathbb{R}^d$
- In this course, f will typically be a neural network



Why learn embeddings?

- Eliminates reliance on handcrafted features
- May capture network information that is not in engineered features
- Potentially useful for many downstream tasks
- Pretraining with self-supervised learning can reduce labeled data requirements for downstream tasks

Embedding in $\mathbb{R}^d$



Pre-trained node or graph
via self-supervised
learning



Downstream Tasks

- Node classification
- Graph classification
- Clustering
- Graph generation
- ... and so on



Example node embeddings

Image credit: Perozzi 2014





Image credit: Perozzi 2014



Embeddings in $\mathbb{R}^2$ for Zachery's Karate network. Notice how community structure of the graph is maintained by the embeddings.



Shallow Node Embedding Framework





Encoder-decoder conceptual description

- We will adopt an encoder-decoder paradigm where:
 - Encoder maps each node (or edge) to an embedding, i.e., a vector in $\mathbb{R}^d$
 - Decoder takes one or more embeddings and infers information about a node, edge, or graph.
 - Typically, we will only care about the encoder for downstream tasks, but both are necessary to form the embeddings




Image credit: Hamilton 2020



Encoder-decoder components

- The Encoder, *ENC*, is a function (e.g., a neural network) that maps nodes to an $\mathbb{R}^d$ embedding
- A node (or edge) similarity measure, $\mathbf{S}[u, v]$, that we'll use to assess performance of the encoder-decoder model
- The Decoder, *DEC*, infers the similarity measure from embedding inputs
- Optimization approach to learn model parameters so that:


$$DEC(\mathbf{e}_u, \mathbf{e}_v) \approx \mathbf{S}[u, v] \text{ for all nodes (or edges) in the network}$$

$$DEC(ENC(u), ENC(v)) \approx \mathbf{S}[u, v] \text{ for all nodes (or edges) in the network}$$



Encoder concept in detail


- The encoder, $ENC(v) = \mathbf{z}_v$, is any function (e.g., a neural network) that maps a node, v , to a vector, $\mathbf{z}_v \in \mathbb{R}^d$
- d is often called the embedding dimension or size
- Today, we will focus on shallow methods that only encode network structure
- Later, we'll see graph neural networks that also use node features





Encoder as a lookup function

- In the simplest case:
 - The graph is static (no nodes added or removed)
 - The embedding is constant for every node and stored in an embedding matrix, $\mathbf{Z} \in \mathbb{R}^{n \times d}$ where d is the embedding dimension and n is the number of nodes
 - The encoder is just a lookup, i.e., given a node ID, i , it returns the i^{th} row from $\mathbf{Z}$





Embedding learning through loss optimization

- We consider a **pairwise decoder**, $DEC(\mathbf{z}_u, \mathbf{z}_v)$, that is a function of any two node embedding inputs
- Let $\mathbf{S}[u, v] \in \mathbb{R}^+$ be a similarity measure of nodes u and v
- Our goal is to construct an encoder-decoder model such that

$$DEC(\mathbf{z}_u, \mathbf{z}_v) \approx \mathbf{S}[u, v]$$

- Standard approach is to minimize the empirical loss, $\mathcal{L}$, over a set of training node pairs $\mathcal{D}$:

$$\mathcal{L} = \sum_{(u,v) \in \mathcal{D}} l[DEC(\mathbf{z}_u, \mathbf{z}_v), \mathbf{S}[u, v]]$$

where $l \in \mathbb{R} \times \mathbb{R} \rightarrow \mathbb{R}$ is a loss function

- We will discuss specific similarity measures and loss functions in the following slides



Shallow node embedding qualities

- Unsupervised or self-supervised learning task (depends on specific method)
- Node features are not used
- Node labels are not used
- Goal is to learn latent features that the decoder can use to reconstruct a similarity measure derived from network structure
- The node embeddings are intended to be useful for any task where network structure is informative



Node Embedding with Matrix Factorization



Matrix factorization embeddings

- Simplest node similarity: Nodes u and v are similar if they are connected by an edge
- Let the decoder be $DEC(\mathbf{z}_u, \mathbf{z}_v) = \mathbf{z}_v^T \mathbf{z}_u$ (inner product decoder)
- This means we want : $\mathbf{z}_v^T \mathbf{z}_u = a_{u,v}$ which is the (u, v) entry of the graph adjacency matrix $\mathbf{A}$
- Equivalently, letting $\mathbf{Z} \in \mathbb{R}^{n \times d}$ be the embedding matrix, we want $\mathbf{Z}\mathbf{Z}^T = \mathbf{A}$




Image credit: Jure Leskovec, <http://cs224w.stanford.edu>

Matrix factorization embeddings

- The embedding dimension d (number of columns in $\mathbf{Z}$) is much smaller than number of nodes n
- Exact factorization $\mathbf{Z}\mathbf{Z}^T = \mathbf{A}$ is generally not possible
- However, we can learn $\mathbf{Z}$ approximately
- Objective function: $\min_{\mathbf{Z}} \|\mathbf{A} - \mathbf{Z}\mathbf{Z}^T\|_2$ - here the similarity function is the adjacency matrix, $\mathbf{S} = \mathbf{A}$
 - We optimize $\mathbf{Z}$ such that it minimizes the L2 norm (Frobenius norm)

$$\|\mathbf{A} - \mathbf{Z}\mathbf{Z}^T\|_2 = \sqrt{\sum_{i=1}^n \sum_{j=1}^n \left(a_{ij} - (\mathbf{z}_i^T \mathbf{z}_j) \right)^2}$$

- Inner product decoder with node similarity defined by edge connectivity is equivalent to matrix factorization of $\mathbf{A}$



Matrix factorization embeddings


- Other methods adopt the same approach for matrix factorization embeddings
 - Same inner product decoder $DEC(\mathbf{z}_u, \mathbf{z}_v) = \mathbf{z}_v^T \mathbf{z}_u$
 - Same objective function $\min_{\mathbf{Z}} \|\mathbf{S} - \mathbf{Z}\mathbf{Z}^T\|_2$
- The difference is how the similarity, $\mathbf{S}$, is defined (there are many alternatives)



Node Embedding with Random Walks



Random walk



Starting at node 7, and taking 5 random steps yields the walk:

$7 \rightarrow 2 \rightarrow 1 \rightarrow 0 \rightarrow 7 \rightarrow 6$

1. Start at a randomly selected node, u . Set $N = n(u)$ to the neighbors of u
2. Randomly pick a node v from N .
3. Set $u = v$
4. Set $N = n(u)$
5. Repeat 2-4 for $T - 1$ times where T is the desired path length (edge count)
 - Loops are allowed, including going back to the previous node from the current node

Random walk embedding concept

- Intuition: Two nodes that co-occur on short random walks are likely to be structurally similar (share neighbors) and should have similar embeddings
- We can express this intuition mathematically by **setting the similarity function, $\mathbf{S}[u, v]$, to**

$$\mathbf{S}[u, v] \equiv p_{G,T}(v|u)$$

where $p_{G,T}(v|u)$ is the probability of visiting node v on a length- T random walk starting at node u on graph G


- Our goal is to develop an encoder-decoder model such that

$$DEC(\mathbf{z}_u, \mathbf{z}_v) \approx p_{G,T}(v|u)$$

for any pair of nodes u and v with embeddings $\mathbf{z}_u$ and $\mathbf{z}_v$

Random walk similarity measure

- In the previous slide, we set $\mathbf{S}[u, v] \equiv p_{G,T}(v|u)$
- But we don't know $p_{G,T}(v|u)$. What can we do?
- Generate many length- T random walks using a walk strategy R
- Estimate probabilities empirically from the random walks



Length-3 Random Walks

$7 \rightarrow 2 \rightarrow 1$
 $7 \rightarrow 0 \rightarrow 1$
 $7 \rightarrow 3 \rightarrow 7$
 $7 \rightarrow 6 \rightarrow 5$
 $7 \rightarrow 0 \rightarrow 1$

Probability Estimates

$p_{G,3}(0|7) = 2/5$
 $p_{G,3}(1|7) = 3/5$
 $p_{G,3}(2|7) = 1/5$
 $p_{G,3}(3|7) = 1/5$
 $p_{G,3}(4|7) = 0/5$
 $p_{G,3}(5|7) = 1/5$
 $p_{G,3}(6|7) = 1/5$



Random walk decoder

- A natural choice for a decoder to approximate $p_{G,T}(v|u)$ is the softmax

$$DEC(\mathbf{z}_u, \mathbf{z}_v) = \frac{\exp(\mathbf{z}_u^T \mathbf{z}_v)}{\sum_{n \in V} \exp(\mathbf{z}_u^T \mathbf{z}_n)}$$

where V is the set of all nodes in the graph

- This will result in all values of $DEC(\mathbf{z}_u, \mathbf{z}_v)$ being in $[0,1]$ and $\sum_{v \in V} DEC(\mathbf{z}_u, \mathbf{z}_v) = 1$ which can be interpreted as a probability distribution
- We want to optimize the encoder-decoder, that is we want to learn embeddings such that

$$DEC(\mathbf{z}_u, \mathbf{z}_v) = \frac{\exp(\mathbf{z}_u^T \mathbf{z}_v)}{\sum_{n \in V} \exp(\mathbf{z}_u^T \mathbf{z}_n)} \approx p_{G,T}(v|u)$$



Naïve random walk embedding optimization

- We could attempt to minimize the square difference

$$\left(\frac{\exp(\mathbf{z}_u^T \mathbf{z}_v)}{\sum_{n \in V} \exp(\mathbf{z}_u^T \mathbf{z}_n)} - \hat{p}_{G,T}(v|u) \right)^2$$

where $\hat{p}_{G,T}(v|u)$ is an empirical estimate of $p_{G,T}(v|u)$ based on random walk samples.

- Intractable for large graphs as requires counting co-occurrences for every pair of nodes which is $O(n^2)$
- In most networks, most pairs will never co-occur on short walks leading to 0 empirical probability which could yield large negative embeddings
- What else can we try?



Random walk embedding optimization by maximum likelihood

- Let's consider a given node u and the set of all nodes $N_R(u)$ visited on some random walk starting at u using a random walk policy R
- Applying maximum likelihood theory, we seek to minimize the objective function

$$\mathcal{L} = \sum_{u \in V} \sum_{v \in N_R(u)} -\log \left[\frac{\exp(\mathbf{z}_u^T \mathbf{z}_v)}{\sum_{n \in V} \exp(\mathbf{z}_u^T \mathbf{z}_n)} \right]$$

Sum over all nodes u

Sum over nodes v seen on random walks starting from u

Estimated probability of u and v co-occurring on a random walk

- Optimization means finding embeddings $\mathbf{z}$ that minimize $\mathcal{L}$



Random walk embedding optimization by maximum likelihood

- Why does this work?

$$\begin{aligned}\mathcal{L} &= \sum_{u \in V} \sum_{v \in N_R(u)} -\log \left[\frac{\exp(\mathbf{z}_u^T \mathbf{z}_v)}{\sum_{n \in V} \exp(\mathbf{z}_u^T \mathbf{z}_n)} \right] \\ &= \sum_{u \in V} \sum_{v \in N_R(u)} -\log[\exp(\mathbf{z}_u^T \mathbf{z}_v)] + \log \left[\sum_{n \in V} \exp(\mathbf{z}_u^T \mathbf{z}_n) \right]\end{aligned}$$

- Minimizing this function requires
 - Pushing $\uparrow \mathbf{z}_u^T \mathbf{z}_v$ - requires $\mathbf{z}_u$ and $\mathbf{z}_v$ to be similar (ideally, correlated)
 - Pushing $\downarrow \sum_{n \in V} \exp(\mathbf{z}_u^T \mathbf{z}_n)$ - requires $\mathbf{z}_u$ to be dissimilar (ideally, anticorrelated) from all $\mathbf{z}_n$ especially $\mathbf{z}_n \notin N_R(u)$
- Eliminates the need to count co-occurrence for every pair of nodes and addresses
0 empirical probability issues, but ...



Random walk embedding optimization by maximum likelihood

- Too computationally expensive

$$\mathcal{L} = \sum_{u \in V} \sum_{v \in N_R(u)} -\log \left[\frac{\exp(\mathbf{z}_u^T \mathbf{z}_v)}{\sum_{n \in V} \exp(\mathbf{z}_u^T \mathbf{z}_n)} \right]$$

Nested summation over all nodes in the graph gives $O(n^2)$ complexity. This is prohibitive on large graphs.



Negative sampling

Revise our loss function using the approximation*

$$-\log \left[\frac{\exp(\mathbf{z}_u^T \mathbf{z}_v)}{\sum_{n \in V} \exp(\mathbf{z}_u^T \mathbf{z}_n)} \right]$$

random distribution
over all nodes

$$\approx -\log[\sigma(\mathbf{z}_u^T \mathbf{z}_v)] - \sum_{i=1}^k \log[\sigma(-\mathbf{z}_u^T \mathbf{z}_{n_i})], n_i \sim P_V$$

sigmoid function

$$\sigma(z) = \frac{1}{1+e^{-z}}$$

- Normalize against k randomly selected *negative samples*
- A form of *Noise Contrastive Estimation (NCE)*
- Eliminates the need to compute the denominator over all nodes
- Complexity is independent of number of nodes in network

*See <https://arxiv.org/pdf/1402.3722>



Negative sampling

Revise our loss function using the approximation*

$$-\log \left[\frac{\exp(\mathbf{z}_u^T \mathbf{z}_v)}{\sum_{n \in V} \exp(\mathbf{z}_u^T \mathbf{z}_n)} \right]$$

$$\approx -\log[\sigma(\mathbf{z}_u^T \mathbf{z}_v)] - \sum_{i=1}^k \log[\sigma(-\mathbf{z}_u^T \mathbf{z}_{n_i})], n_i \sim P_V$$

- Sample k negative nodes n_i with probability proportional to node degree
- Higher k gives more robust estimates but also higher bias on negative events (meaning, the negative sample node might actually occur on random walks we just didn't take enough samples)
- In practice, $k \in [5, 20]$




*See <https://arxiv.org/pdf/1402.3722>



NCE (loss function version)

$$-\left[\log[\sigma(\mathbf{z}_u^T \mathbf{z}_v)] + \sum_{i=1}^k \log[\sigma(-\mathbf{z}_u^T \mathbf{z}_{n_i})], n_i \sim P_V\right]$$

Contour plot of the NCE loss function. The x-axis is the value of $\mathbf{z}_u^T \mathbf{z}_v$. The y-axis is the value of $\mathbf{z}_u^T \mathbf{z}_{n_i}$.




The orange curve represents the value of $-\log[\sigma(\mathbf{x})]$ where $\mathbf{x} = \mathbf{z}_u^T \mathbf{z}_v$. The best we can do is make this dot product large. The green curve represents $-\log[\sigma(-\mathbf{x})]$ where $\mathbf{x} = \mathbf{z}_u^T \mathbf{z}_{n_i}$. The best we can do is make this value negative with large magnitude.



Understanding NCE loss

$$-\left[\log[\sigma(\mathbf{z}_u^T \mathbf{z}_v)] + \sum_{i=1}^k \log[\sigma(-\mathbf{z}_u^T \mathbf{z}_{n_i})], n_i \sim P_V\right]$$



Key Insight: Both similar vectors (positive pairs) and anti-correlated vectors (negative pairs) achieve low loss values. The model learns both *attraction* and *repulsion* relationships!

A. Similar Vectors (Positive Pair)

$$\mathbf{z}_1 = [3, 4], \mathbf{z}_{\text{pos}} = [4.5, 6]$$

$$\text{dot_product} = 3 \times 4.5 + 4 \times 6 = 37.5$$

$$\text{sigmoid}(37.5) \approx 1.000$$

$$\log(\text{sigmoid}(37.5)) \approx 0.000$$

✓ Low loss for positive pair - vectors learn to align

B. Anti-correlated Vectors (Negative Pair)

$$\mathbf{z}_1 = [3, 4], \mathbf{z}_{\text{neg}} = [-6, -8]$$

$$\text{dot_product} = 3 \times (-6) + 4 \times (-8) = -50$$

$$\text{sigmoid}(-(-50)) = \text{sigmoid}(50) \approx 1.000$$

$$\log(\text{sigmoid}(50)) \approx 0.000$$

✓ Low loss for negative pair - strong opposition is good!

C. Orthogonal Vectors (Negative Pair)

$$\mathbf{z}_1 = [3, 4], \mathbf{z}_{\text{neg}} = [4, -3]$$

$$\text{dot_product} = 3 \times 4 + 4 \times (-3) = 0$$

$$\text{sigmoid}(-0) = \text{sigmoid}(0) = 0.500$$

$$\log(\text{sigmoid}(0)) \approx -0.693$$

↑ Higher loss - neutral relationship, room for improvement



How do we optimize our objective function?

- Assume our loss function is $\mathcal{L}$ and we want to optimize the embedding $\mathbf{z}_u$
- Stochastic Gradient Descent – **process one sample at a time**
 1. Randomly initialize $\mathbf{z}_m$ for all nodes m
 2. For each random walk sample $s = \{u, v_1, \dots, v_w, n_1, \dots, n_k\}$
 - a. Compute the partial derivatives $\frac{\delta \mathcal{L}}{\delta \mathbf{z}_m}$ for all nodes $m \in s$
 - b. For all nodes m , step in the negative gradient direction $\mathbf{z}_m \leftarrow \mathbf{z}_m - \alpha \frac{\delta \mathcal{L}}{\delta \mathbf{z}_m}$ where α is the learning rate
 3. Repeat 2 until convergence
- Completion of step 2 represents one epoch



Random walk summary


1. Generate short fixed-length random walks starting from each node, u , on the graph using a random walk policy R
2. For each node u , determine the neighborhood $N_R(u) = \{v_1, \dots, v_w\}$ from the walks and combine with k randomly selected negative nodes $\{n_1, \dots, n_k\}$. This represents a single sample $s = \{u, v_1, \dots, v_w, n_1, \dots, n_k\}$
3. Optimize the embeddings $\mathbf{Z}$ using stochastic gradient descent with the NCE loss function (using negative sampling)

$$\mathcal{L}_s = \sum_{v \in N_R(u)} -\log[\sigma(\mathbf{z}_u^T \mathbf{z}_v)] - \sum_{i=1}^k \log[\sigma(-\mathbf{z}_u^T \mathbf{z}_{n_i})]$$



What is a random walk policy?


- As of yet, we haven't specified the random walk policy R
- Random walk policy definition
 - Assume the random walk is at node v
 - *Set a rule (policy) for deciding which neighbor node of v to select for the next step*
 - The rule can depend on parameters and other information about the walk state and network





How do we pick the random walk policy?

- Simplest approach is to select the next node in the walk from the neighbors of the current node with equal probability
 - Unbiased random walk
 - This was proposed for [DeepWalk](#) by Perozzi et al.
 - **This turns out to be too restrictive**
- How can we improve this?



Unbiased Random Walk Policy Example

Node v has 5 neighbors. Pick a number, ρ , from a uniform random distribution on $[0,1]$.

Pick the next node, u , with the rule:

$\rho \in [0, 0.2) \rightarrow u = 1$

$\rho \in (0.2, 0.4) \rightarrow u = 2$

$\rho \in (0.4, 0.6) \rightarrow u = 3$

$\rho \in (0.6, 0.8) \rightarrow u = 4$

$\rho \in (0.8, 1] \rightarrow u = 5$



Breadth and depth first search

- **Breadth-first** is a walk policy that favors a local neighborhood view
- **Depth-first** is a walk policy that favors a global neighborhood views
- We would like a walk policy that combines these
- Next, we discuss *node2vec* which uses a parameterized bias random walk policy that interpolates between **breadth-first (local)** and **depth-first (global)** views of a node neighborhood





Image credit: Grover & Leskovec 2016



node2vec uses a biased random walk policy

- Biased random walk concept
 - Assume the random walk just traversed the edge from node t to node v
 - **Bias the probability of picking a neighbor of node v based on whether it is**
 - The previous node
 - A neighbor of the previous node
 - Not a neighbor of the previous node
 - How do we bias the probability?






node2vec uses a biased random walk policy

- Biased probabilities
 - Assume the random walk just traversed the edge from node t to node v
 - Select policy parameters p and q (more on these later)
 - For node x in set of neighbors of v
 - Let w_{vx} be the weight on the edge between nodes v and x (unweighted $w_{vx} = 1$)
 - Let d_{tx} be the shortest path distance from node t to node x (d_{tx} is 0, 1, or 2)
 - Set the unnormalized transition probability to $\pi_{vx} = \alpha_{pq}(t, x) \cdot w_{vx}$
 - Normalize the transition probabilities by $\sum_{x \in N(v)} \pi_{vx}$

$$\alpha_{pq}(t, x) = \begin{cases} \frac{1}{p} & \text{if } d_{tx} = 0 \\ 1 & \text{if } d_{tx} = 1 \\ \frac{1}{q} & \text{if } d_{tx} = 2 \end{cases}$$





node2vec biased random walk step example

- Let $\Pi_v = \sum_{x \in N(v)} \pi_{vx}$ be the sum of the unnormalized transition probabilities $\pi_{vx} = \alpha_{pq}(t, x) \cdot w_{vx}$ so that the normalized transition probability from node v to x is $\psi_{vx} = \pi_{vx} / \Pi_v$.

Example


Let $w_{vx} = 1$ for all x (unweighted)

Let $p = 0.5$ and $q = 0.25$

$\pi_{vt} = 2, \pi_{v1} = \pi_{v2} = 1, \pi_{v3} = \pi_{v4} = 4$

$$\Pi_v = \sum_{x \in N(v)} \pi_{vx} = 2 + 2 \cdot 1 + 2 \cdot 4 = 12$$

$$\alpha_{pq}(t, x) = \begin{cases} \frac{1}{p} & \text{if } d_{tx} = 0 \\ 1 & \text{if } d_{tx} = 1 \\ \frac{1}{q} & \text{if } d_{tx} = 2 \end{cases}$$



Pick a number, ρ , from a uniform random distribution on $[0,1]$. Pick the next node, u , with the rule:

$\rho \in [0, 1/12) \rightarrow u = 1$
 $\rho \in (1/12, 2/12) \rightarrow u = 2$
 $\rho \in (2/12, 4/12) \rightarrow u = t$
 $\rho \in (4/12, 8/12) \rightarrow u = 3$
 $\rho \in (8/12, 1] \rightarrow u = 4$



node2vec uses a biased random walk policy

- Random walk policy parameters
 - **In-out parameter q** : scales the probability of staying **in (BFS)** the previous node's neighbors vs. moving **out (DFS)** of the previous node's neighbors (DFS)
 - Setting $q > 1$ encourages BFS and $q < 1$ encourages DFS (we saw this in the previous example)
 - **Return parameter p** : scales the probability of returning to the previous node in a walk
 - Setting $p > \max(q, 1)$ discourages backtracking, setting $p < \min(q, 1)$ encourages backtracking, otherwise $\min(q, 1) < p < \max(q, 1)$ has moderate backtracking



node2vec context sampling


- Assume we have created a large number of random walks with the node2vec biased walk policy
- How do we use them to create samples?
- Consider a single random walk
 1. Set the context window length w
 2. Pick a center node, u , in the walk
 3. Let every node, v , that is within w steps to u be a context node
 4. Form positive pairs (u, v) for u and every context node
 5. Form a single sample as the set of positive pairs and k randomly sampled negative nodes (i.e., those not in the graph that are not in the context nodes)



node2vec context sampling



A single random walk
of length 6




Context Window for center node n7 at position 2

Context nodes: n1, n3, n12, n3

Window range: [0, 4]

Positive pairs:

(n7, n1), (n7, n3), (n7, n12), (n7, n3)



Context Window for center node n1 at position 5

Context nodes: n12, n3, n23


Window range: [3, 6]

Positive pairs:

(n1, n12), (n1, n3), (n1, n23)



node2vec context sampling



Each row is a single sample for SGD

Sample	Walk	Center Pos	Center Node	Context	Positive Pairs	Negative Samples
1	1	0	n1	n3, n7	(n1, n3), (n1, n7)	n45, n29, n86, n63, n24
2	1	1	n3	n1, n7, n12	(n3, n1), (n3, n7), (n3, n12)	n41, n86, n99, n72, n27
3	1	2	n7	n1, n3, n12, n3	(n7, n1), (n7, n3), (n7, n12), (n7, n3)	n14, n5, n81, n61, n66
4	1	3	n12	n3, n7, n3, n1	(n12, n3), (n12, n7), (n12, n3), (n12, n1)	n44, n68, n15, n81, n37
5	1	4	n3	n7, n12, n1, n23	(n3, n7), (n3, n12), (n3, n1), (n3, n23)	n74, n15, n28, n20, n19
6	1	5	n1	n12, n3, n23	(n1, n12), (n1, n3), (n1, n23)	n81, n2, n54, n87, n26
7	1	6	n23	n3, n1	(n23, n3), (n23, n1)	n83, n29, n27, n45, n87



node2vec algorithm summary

1. Precompute edge transition probabilities for all edges
2. Generate r random walks of length l for every node in the graph using the bias random walk policy
3. Use context sampling to generate samples where each sample include positive pairs (u, v) and k negative nodes
4. Apply SGD over each sample using NCE with negative sampling loss
 - Linear time complexity
 - All steps are individually parallelizable



Random walk variants

- Different kinds of biased random walks:
 - Based on node attributes ([Dong et al.](#))
 - Based on learned weights ([Abu-El-Haija et al.](#))
- Alternative optimization schemes:
 - Directly optimize based on 1-hop and 2-hop random walk probabilities (as in [LINE from Tang et al.](#))
- Network preprocessing techniques:
 - Run random walks on modified versions of the original network (e.g., [Ribeiro et al. struct2vec](#), [Chen et al. HARP](#))



Summary so far

- **Core idea:** Embed nodes so that distances in embedding space reflect node similarities in the original network
- **Different notions of node similarity:**
 - Naïve: Similar if two nodes are connected
 - Random walk approaches



Summary so far

- So, what method should I use..?
- **No one method wins in all cases....**
 - E.g., node2vec performs better on node classification while alternative methods perform better on link prediction ([Goyal and Ferrara, survey](#))
- Random walk approaches are generally more efficient
- In general: Must choose definition of node similarity that matches your application



Graph Embeddings






Embedding an entire graph

Goal: Embed a subgraph (or the entire graph) G in the embedding space

Motivation: Classify or cluster networks


Examples: candidate drug targets, anomalous networks






Node averaging for graph embedding

- Apply a node embedding algorithm to all nodes in G
- Average the node embeddings to obtain a single embedding for G
- The average is done element wise




position i in G embedding is the average of the embedding values at position i from the node embeddings





Virtual nodes for graph embedding

- Create a virtual node and connect all nodes in the subgraph (or entire graph) to the virtual node
- Run a node embedding method to obtain the embedding for the virtual node





Limitations of matrix factorization and random walk embeddings





Matrix factorization and random walk embedding methods are transductive

- Transductive (not inductive) method:
 - Cannot obtain embeddings for nodes not in the training set
 - Cannot apply to new graphs using the same nodes or the same graph with updated edges
 - If you apply a random walk method to the same graph multiple times, you'll get a different embedding each time. Why?




Image credit: Jure Leskovec, <http://cs224w.stanford.edu>



Random walk embedding methods cannot capture structural similarity

- DeepWalk and node2vec do not capture structural similarity
- Consider the network below
 - Node 1 and 11 are structurally similar – part of one triangle, degree 2, ...
 - However, they have very different random walk embeddings.
 - Why? It's unlikely that a random walk will reach node 11 from node 1





Image credit: Jure Leskovec, <http://cs224w.stanford.edu>



Matrix factorization and random walk embedding methods cannot use node, edge, or graph features



Node feature vector

(e.g. protein properties in a protein-protein interaction graph)

Matrix factorization / random walk embeddings do not incorporate such node features

Solution: Graph neural networks

Image credit: Jure Leskovec, <http://cs224w.stanford.edu>



What can we do with node embeddings?

- Clustering/community detection: Cluster points $\mathbf{z}_n$
- Node classification: Predict label of node n based on $\mathbf{z}_n$
- Link prediction: Predict edge (u, v) based on a function of $f(g(\mathbf{z}_u, \mathbf{z}_v))$ where the function g is usually one of vector concatenation, averaging, Hadamard product, or difference
- Graph classification: Predict label based on graph embedding, $\mathbf{z}_G$, formed via aggregating node embeddings or virtual-node

